The AI Investment Decision Framework

Should You Actually Use AI for This?

I've advised over 50 SMBs on AI adoption. I've watched companies burn $50,000 on custom LLM deployments that resulted in zero efficiency gains, and I've seen a $500 weekend experiment with Zapier and Claude save a logistics company $80,000 a year in manual data entry.

The difference is never the technology. It's the target.

Most businesses look at AI and ask, "What can this do?" That is a guaranteed way to waste money. You should be asking, "What painful, expensive, repetitive bottleneck do we already have that a machine can execute?"

This framework is designed to force hard thinking. No 1-10 wishy-washy scales. No buzzwords. Just a brutal filter to ensure you only invest in AI where it will actually work.

Prerequisite: The 6-Month Rule

If you haven't done the task manually for 6 months, you cannot automate it.

Why: AI is not a magic wand for broken processes. It is a strict amplifier. If you apply AI to a disorganized, undefined process, you will just get messy, unpredictable results at the speed of light. Before you bring in an LLM, a human needs to have bled over the task long enough to define the edge cases, the standard operating procedures, and what "good" looks like.

If you are launching a brand new service line today, do it manually. Revisit AI in six months.

The 4-Question Scoring Framework

To proceed with an AI project, the task must score four hard "YES" answers. One "NO" means the project dies or requires a fundamental redesign.

Question 1: TIME

"Does this consume >3 hours/week of someone's actual time?"

- **The Reality Check:** Humans are terrible at estimating time. Before answering, make the employee run a Toggl or Clockify timer on this specific task for one week.
- **The Trap:** People often want to automate tasks they *hate* doing, even if it only takes 15 minutes a week. Setting up, testing, and maintaining an AI workflow for a 15-minute task will take more time than just doing the task manually for the next five years.
- **RED FLAG:** "It only takes 20 minutes but it's really annoying." Kill the project.

Question 2: REPEATABILITY

"Is the task 80%+ identical each time?"

- **The Reality Check:** AI excels at patterned execution, not novel strategy. Can you write a bulletproof, 10-step checklist for this task that a smart 19-year-old intern could follow perfectly without asking questions? If yes, AI can do it.
- **The Creative Work Trap:** Don't use AI to invent your go-to-market strategy or design a brand-new product from scratch. Use it to extract 50 invoice numbers from 50 wildly different PDF formats and put them in a CSV. High variance in input format is fine; high variance in the required *logic* is fatal.
- **RED FLAG:** "Every client situation is completely unique and requires a custom approach." Kill the project.

Question 3: ERROR COST

"If the AI gets it wrong, can you catch it in <10 minutes and fix it for <$50?"

- **The Reality Check:** AI hallucinates. It will confidently lie to you. Your workflow must assume a 5% failure rate. Therefore, there must be a "human in the loop" to catch the garbage before it hits the customer or the bank account.
- **War Stories:** * *The $1 Chevy:* A car dealership deployed an unmonitored AI chatbot that agreed to sell a customer a 2024 Chevy Tahoe for $1. (Legal nightmare).
    - *The Phantom Precedent:* A lawyer used ChatGPT to write a brief without checking the citations. ChatGPT invented fake court cases. The lawyer was sanctioned.
- **RED FLAG:** "The output goes directly to the client/live database without human review." Kill the project.

Question 4: DATA SAFETY

"Can you do this without feeding proprietary customer, employee, or financial data into a public model?"

- **The Reality Check:** If you paste your client's raw P&L into the free, consumer version of ChatGPT, you just violated their NDA. You are feeding their data into OpenAI's training apparatus.
- **RED FLAG:** "We'll just copy-paste the customer database into the chat window."

The 1-Page Vendor Security Assessment (If using a third-party AI tool):

Before buying an AI wrapper or SaaS tool, email their sales rep this exact list. If they dodge, walk away.

1. **Zero Data Retention / Non-Training:** "Do you explicitly state in your enterprise MSA that our data is NEVER used to train your foundational models?" (Must be YES).
2. **Data Residency:** "Where exactly is our data hosted and processed?"
3. **Compliance:** "Are you SOC2 Type II certified?"
4. **Third-Party APIs:** "Are you passing our data to OpenAI/Anthropic? If yes, are you doing it via zero-retention Enterprise APIs?"

Tool Selection Matrix

Stop arguing about which model is "best." Use the right tool for the job.

Tool	When to Use It	When to Avoid It
Claude (Anthropic)	Writing that sounds human, coding, and reading massive documents. It has the best "context window" and follows complex formatting instructions brilliantly.	Live web browsing or deep ecosystem integrations.
ChatGPT (OpenAI)	General reasoning, data analysis (Advanced Data Analysis is elite), voice conversations, and custom GPTs for internal team use.	Writing copy (it naturally sounds pompous and robotic).
Gemini (Google)	You are already a Google Workspace power user. Summarizing long YouTube videos, pulling data out of Gmail/Drive instantly.	Complex coding tasks (Claude is currently better).
Custom API / Make.com	You passed the 4-question framework and need this process to run in the background 100 times a day automatically.	You haven't proven the prompt works manually yet.

The "Kill List": 5 AI Projects That Always Fail

If someone pitches one of these, fire them or throw them out of the room.

1. **Fully Autonomous Customer Support for High-Ticket Items:** If a client pays you $10,000/month, they do not want to talk to an AI widget when something breaks. They want a human throat to choke.
2. **"The Everything Chatbot":** "Let's connect an AI to our entire company Google Drive so employees can ask it anything!" It will hallucinate, pull up outdated HR policies from 2018, and accidentally expose executive payroll data. Start with a tightly scoped, single-purpose database.
3. **The Unsupervised Sales Outreach Cannon:** Hooking AI up to LinkedIn/Email to send 1,000 hyper-personalized messages a day. The AI will inevitably misinterpret a prospect's bio and send an incredibly offensive or nonsensical email. It burns your domain reputation.
4. **Replacing the Junior Copywriter Entirely:** AI writes structure well; it writes soul terribly. If you fire the writer and just use AI, your brand will sound like B2B SaaS cardboard in three months. Use AI to generate 50 bad ideas; pay the human to find the 1 good one.
5. **Complex Multi-Agent Workflows for Simple Forms:** Don't build an ecosystem of 4 AI agents talking to each other to qualify a lead. Just use a Typeform.

The AI Project Scoring Sheet (Printable)

Project Name: _____

Proposed By: _____

Date: _____

Prerequisite Check

[ ] **The 6-Month Rule:** We have executed this exact process manually for at least 6 months. We know the edge cases.

If unchecked, STOP. Project is killed.

The 4-Question Framework

(All questions must be answered YES to proceed).

1. TIME: Does this task consume >3 hours of actual, tracked time per week?

[ ] YES

[ ] NO *(Red Flag: Waste of setup time)*

2. REPEATABILITY: Is the underlying logic of this task 80%+ identical every single time?

[ ] YES

[ ] NO *(Red Flag: Requires novel, strategic human thought)*

3. ERROR COST: If the AI hallucinates, can a human catch it in <10 minutes and fix it for <$50 without client impact?

[ ] YES

[ ] NO *(Red Flag: Unsupervised outputs or catastrophic liability)*

4. DATA SAFETY: Can we execute this without putting proprietary/PII data into a public model that trains on our inputs?

[ ] YES

[ ] NO *(Red Flag: Massive security/NDA violation)*

Final Verdict

Score (Count of YES answers): ____ / 4

- **4/4 YES:** Approved for prototyping.
- **Anything less than 4:** KILLED. Back to the drawing board.

Next Step (If Approved):

Tool to be used for initial prototype: _____

Name of the "Human in the Loop" who will review outputs: _____